

Moulshree Guleria

Chicago, IL - 60607 | mxguleria@gmail.com | 312-687-5439 | LinkedIn

Education

- University of Illinois at Chicago**, MS in Computer Science Aug 2024 – May 2026
- GPA: 4.0/4.0; Teaching Assistant for CS 411 - Artificial Intelligence I (Spring 2025 and Fall 2025)
 - **Coursework:** Artificial Intelligence, Engineering Distributed Objects for Cloud Computing, Neural Engineering

Nitte Meenakshi Institute of Technology, BE in Information Science Aug 2016 – Aug 2020

 - GPA: 3.615/4.0; Coursework: Data Structures, Algorithms, OOPs, Software Engineering, DBMS

Experience

- Artificial Intelligence Strategy Intern**, The World Music Foundation – Chicago, IL Sept 2025 – Dec 2025
- Engineered AI-powered validation pipeline using Open AI's ChatGPT for automating job applications screening, cutting manual effort by 70%.
 - Collaborated with leadership for AI adoption strategy aligned with compliance, reliance, and long-term goals.
 - Developed WMF website on WordPress with Javascript, HTML, CSS, improving UX and mobile responsiveness.

Senior Software Engineer, PriceWaterhouseCoopers (PwC) – Bengaluru, KA Jan 2021 – Aug 2024

 - Developed and deployed 10+ REST APIs using C# and ASP .NET, on PwC's Talent Exchange.
 - Led an 8-member incident response team resolving operational incidents with 95% on-time efficiency, providing 24/7 on-call support.
 - Served as technical lead translating business requirements into solutions for leadership & financial stakeholders.
 - Streamlined CI/CD pipelines using Azure DevOps, achieving 98% deployment success rate.

Software Engineer Intern, Microexcel Inc. – Mysore, KA June 2019 – July 2019

 - Designed and developed a full stack website using Python and Django with data visualization, weather integration, Oxford dictionary lookup, and GitHub integration.
 - Engineered a face recognition-based attendance system using Python, OpenCV and Django.

Projects

- Distributed GraphRAG System on AWS (UIC)**
- Architected end-to-end RAG system processing 600+ MSR papers through progressive distributed frameworks: Hadoop MapReduce on EMR (85% processing time reduction), Apache Spark with Delta Lake (70% cost reduction via content-hash deduplication, <2s query latency), and Apache Flink on EKS (real-time Neo4j knowledge graph construction with Ollama LLM). Deployed production Kubernetes infrastructure with IAM security, EBS persistence, RESTful APIs (Akka HTTP), and sub-second P95 query latencies. Tools: Hadoop MapReduce, Apache Spark, Apache Flink, AWS (EMR, EKS, S3), Scala, Delta Lake, Neo4j, Kubernetes, Docker, Ollama, Lucene HNSW, Akka HTTP

Memory Anchor – Agentic AI App for Alzheimer's Support
(Google Cloud × ODSC Agentic AI Hackathon Finalist, Top 6/80)

 - Designed and implemented an Agentic AI app that provides routine reminders, memory cues, and context-aware emotional support for early-stage Alzheimer's patients. Built memory and execution modules with Google Gemini and its API to handle personalized conversations and alerts. Integrated onboarding, memory storage, and emergency notification workflows to caregivers. Tools: Python, Google Gemini API, Agentic AI, Email Automation

Movie Recommendation System (UIC)

 - Developed and deployed an AI-powered movie recommendation system achieving 0.92 RMSE using SVD++ algorithms. Built a scalable FastAPI inference API containerized with Docker, serving recommendations in <200 ms. Created an end-to-end ML pipeline with automated retraining and evaluation via Jenkins. Tools: Python, Surprise, FastAPI, Docker, Jenkins, LaunchDarkly, Prometheus, Grafana

Technologies

Cloud: AWS (EMR, EKS, S3, EC2, ECR, IAM), Azure, Azure DevOps, Kubernetes, Docker, Jenkins
Languages: Python, C#, Scala
Web: JavaScript, HTML, CSS3, WordPress
AI/LLM Toolkits: Google Gemini API, OpenAI, Ollama, Agentic AI, RAG, Vector Search (Lucene HNSW)
Big Data & Streaming: Hadoop MapReduce, Apache Spark, Apache Flink, Delta Lake
Databases: Neo4j, SQL, NoSQL, Vector Databases
Frameworks & Backends: ASP.NET, FastAPI, Django, REST APIs, Akka HTTP, Microservices